Otto-Friedrich-University of Bamberg

Professorship for Computer Science,
Communication Services, Telecommunication
Systems and Computer Networks

Foundation of Internet Communication

Assignment 6
Interdomain Routing Protocols

Submitted by:
Group X

Moktahid Al Faisal
Abdullah Al Mosabbir
Mohammed Mehedi Hasan
Kazi Sayef Shawgat
Sheikh Jumon Ahmed

Supervisor: Prof. Dr. Udo Krieger

Bamberg, July 19, 2020
Summer Term 2020

Contents

List of Figures

1 Exterior Routing with BGP

1. Build the topology depicted in Figure 1 with Kathara

   Soluation: Done

2. Determine the addresses of the Internet DNS root servers

   Soluation: Done

3. Configure the BGP peering between both ASes, where AS 1 is responsiblefortheOSPFnetworkandAS2forRIP

   Solution: To complete the task, we set 'bgpd = yes' in the daemons files of core3 and rip4. And also create a bgpd.conf file for both of the routers.

```
router bgp 1
!
neighbor 13.0.0.4 remote-as 2
neighbor 13.0.0.4 description Router 2 of AS2
```

Figure 1: bgpd.cong of core3

```
router bgp 2
!
neighbor 13.0.0.3 remote-as 1
neighbor 13.0.0.3 description CORE3 of AS1
```

Figure 2: bgpd.cong of rip4

4. Adjust the announcements, so that all hosts of both ASes can reach each other

   Solution: We have modified the bgpd.conf files like below:

```
router bgp 1
!
network 10.1.1.0/28
network 11.0.0.0/12
network 12.0.0.0/20
!
neighbor 13.0.0.4 remote-as 2
neighbor 13.0.0.4 description Router 2 of AS2
```

Figure 3: bgpd.cong of core3

```
router bgp 2
!
network 50.0.0.0/9
!
neighbor 13.0.0.3 remote-as 1
neighbor 13.0.0.3 description CORE3 of AS1
```

Figure 4: bgpd.cong of rip4

5. Add static routes wherever needed to ensure global connectivity

   Solution: To do this, we have added the command redistribute bgp in 'ospfd.conf ' of core3 and in ripd.conf of rip4.

6. Inspect the routing tables of rip4 and core3 and explain, how the peering is done with BGP

   Solution: In both routers core3 and rip4, bgp sends the corresponding connected networks to its neighbor ASes. We can see the updated routing tables below.

```
core3 [/]# ip route show
10.1.1.0/30 dev eth0 proto kernel scope link src 10.1.1.2
10.1.1.4/30 dev eth1 proto kernel scope link src 10.1.1.6
10.1.1.8/30 nhid 26 proto ospf metric 20
        nexthop via 10.1.1.5 dev eth1 weight 1
        nexthop via 10.1.1.1 dev eth0 weight 1
11.0.0.0/15 nhid 27 via 10.1.1.1 dev eth0 proto ospf metric 20
11.2.0.0/16 nhid 27 via 10.1.1.1 dev eth0 proto ospf metric 20
11.4.0.0/18 nhid 27 via 10.1.1.1 dev eth0 proto ospf metric 20
11.4.64.0/28 nhid 27 via 10.1.1.1 dev eth0 proto ospf metric 20
11.4.64.0/20 nhid 27 via 10.1.1.1 dev eth0 proto ospf metric 20
11.8.0.0/13 nhid 27 via 10.1.1.1 dev eth0 proto ospf metric 20
12.0.0.0/23 nhid 16 via 10.1.1.5 dev eth1 proto ospf metric 20
12.0.2.0/23 nhid 16 via 10.1.1.5 dev eth1 proto ospf metric 20
12.0.4.0/25 nhid 16 via 10.1.1.5 dev eth1 proto ospf metric 20
12.0.5.0/24 nhid 16 via 10.1.1.5 dev eth1 proto ospf metric 20
12.0.8.0/21 nhid 16 via 10.1.1.5 dev eth1 proto ospf metric 20
13.0.0.0/20 dev eth2 proto kernel scope link src 13.0.0.3
50.0.0.0/9 nhid 14 via 13.0.0.4 dev eth2 proto bgp metric 20
core3 [/]#
```

Figure 5: Routing table of core3

```
rip4 [/]# ip route show
10.1.1.0/28 nhid 13 via 13.0.0.3 dev eth0 proto bgp metric 20
11.0.0.0/12 nhid 13 via 13.0.0.3 dev eth0 proto bgp metric 20
12.0.0.0/20 nhid 13 via 13.0.0.3 dev eth0 proto bgp metric 20
13.0.0.0/20 dev eth0 proto kernel scope link src 13.0.0.4
50.0.0.0/20 nhid 10 via 50.0.16.5 dev eth1 proto rip metric 20
50.0.0.0/8 dev eth2 proto kernel scope link src 50.0.0.4
50.0.16.0/22 dev eth1 proto kernel scope link src 50.0.16.4
50.1.0.0/24 nhid 10 via 50.0.16.5 dev eth1 proto rip metric 20
50.2.0.0/23 nhid 8 via 50.0.0.3 dev eth2 proto rip metric 20
50.2.2.0/24 nhid 8 via 50.0.0.3 dev eth2 proto rip metric 20
50.3.0.0/18 nhid 8 via 50.0.0.3 dev eth2 proto rip metric 20
50.10.0.0/16 nhid 10 via 50.0.16.5 dev eth1 proto rip metric 20
50.48.0.0/12 nhid 8 via 50.0.0.3 dev eth2 proto rip metric 20
50.64.0.0/10 nhid 8 via 50.0.0.3 dev eth2 proto rip metric 20
rip4 [/]#
```

Figure 6: "Routing table of rip4

2 Multihoming and Redundancy

1. Add AS3 with a new BGP router as3

   Soluation: Done as asked. Added the following commands in the lab.conf file

```
as3[0]="U"
as3[1]="V"
as3[2]="W"
as3[image]="unibaktr/alpine:frr"
as3[exec]="ifconfig eth0 13.0.16.2 netmask 255.255.248.0 up"
as3[exec]="ifconfig eth1 100.0.0.2 netmask 255.255.240.0 up"
as3[exec]="ifconfig eth2 13.0.24.2 netmask 255.255.252.0 up"
```

Figure 7: configurationf for as3 router

2. Attach kili over 100.0.0.0/20 to it. Choose the IP addresses accordingly

   Solution: Added the following commands in the lab.conf file

```
kili[0]="V"
kili[image]="unibaktr/alpine:whoami"
kili[exec]="ifconfig eth0 100.0.0.1 netmask 255.255.240.0 up"
kili[exec]="ip route add default via 100.0.0.2
```

Figure 8: configuration for kili

3. Create CD U to connect as3 with core1 over the network 13.0.16.0/21

4. In the same fashion, create CD W to interconnect as3 with rip1 over the network 13.0.24.0/22. *ru*

   Solution: Done as asked. See the Figure 7

5. Adjust the BGP configuration in a way, that eBGP (external BGP) is used for CDs I, U and W

```
router bgp 3
network 100.0.0.0/20
!
neighbor 13.0.16.1 remote-as 1
neighbor 13.0.16.1 description CORE1 of AS1
neighbor 13.0.16.1 default-originate
!
neighbor 13.0.24.1 remote-as 2
neighbor 13.0.24.1 description RIP1 of AS2
neighbor 13.0.24.1 default-originate
```

Figure 9: bgpd.conf configuration for router as3

```
router bgp 1
network 10.1.1.0/28
network 11.0.0.0/12
network 12.0.0.0/20
network 13.0.16.0/21
!
neighbor 13.0.16.2 remote-as 3
neighbor 13.0.16.2 description as3 of AS3
neighbor 13.0.16.2 default-originate
!
neighbor 10.1.1.2 remote-as 1
neighbor 10.1.1.2 description core3 of AS1 (iBGP)
```

Figure 10: bgpd.conf configuration for router core1

```
router bgp 2
network 50.0.0.0/9
network 13.0.24.0/22
!
neighbor 13.0.24.2 remote-as 3
neighbor 13.0.24.2 description as3 of AS3
neighbor 13.0.24.2 default-originate
```

Figure 11: bgpd.conf configuration for router rip1

6. Additionally, use iBGP (internal BGP) to propagate learned routes inside the networks of AS 1 and AS 2 *ru*

   Soluation: Done

7. Ensure connectivity of all hosts.

   Soluation: Done

8. Now, let us start a short evaluation on the path from bombur to balin.

   1) Determine the path between both nodes with traceroute.:

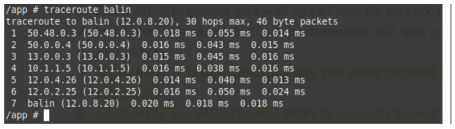

Figure 12: Path from bombur to balin

2) Start a Wireshark capture on the involved CD I or W and remember if rip1 or rip4 are traversed

Solution: Started wireshark capture on CD I, as the traversed path includes router rip4.

3) Start to continuously ping balin from bombur

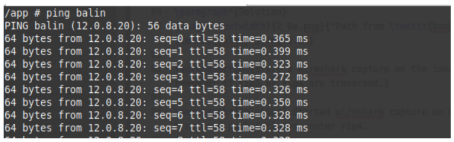

Figure 13: Path from bombur to balin

4) Open a vtysh terminal on the involved RIP router and temporarily remove the BGP neighbor

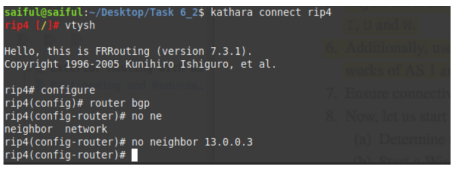

Figure 14: Removing BGP neighbor from rip4

5) Wait until the ttl value of the ping changes

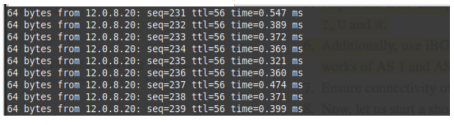

Figure 15: ttl changes after removing BGP neighbor

6) Enable the BGP neighbor again, wait for the ping to change and stop the Wireshark capture

I. Explain BGP's update mechanism based on the Wireshark capture. Solution: It can be seen from the wireshark capture, that the rip4 and core3 exchanges 'KEEPALIVE Message' between them to maintain connecivity. After removing BGP neighbor from rip4, it sends core3 a 'NOTIFICATION Message'. Then they keep exchanging 'OPEN Message'. After reattaching the neighbor, the rip4 sends 'UPDATE Message' to core3. Then core3 updates its routing information and returns an updated 'UPDATE Message' to rip4, after which rip4 does the same thing again.

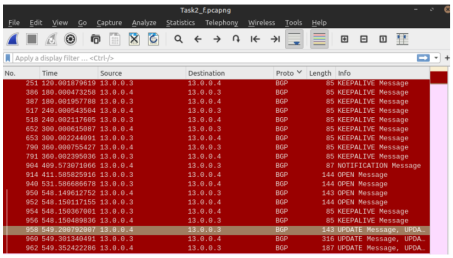

Figure 16: Wireshark capture on CD I

References